

#Music Genre Classification

```
#importing all essential libraries
import random
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#Description of dataset

- 1.filename: Name of the audio file associated with each observation (e.g., blues.00000.0.wav).
- 2.length: Duration of the audio sample in milliseconds.
- 3.chroma_stft_mean and chroma_stft_var: Mean and variance of the chroma short-time Fourier transform (STFT).
- 4.rms_mean and rms_var: Mean and variance of the Root Mean Square (RMS) energy.
- 5.spectral_centroid_mean and spectral_centroid_var: Mean and variance of the spectral centroid.
- 6.spectral_bandwidth_mean and spectral_bandwidth_var: Mean and variance of spectral bandwidth.
- 7.rolloff_mean and rolloff_var: Mean and variance of spectral roll-off frequency.
- 8.zero_crossing_rate_mean and zero_crossing_rate_var: Mean and variance of zero-crossing rate.
- 9.harmony_mean and harmony_var: Mean and variance of harmony.
- 10.percepctr_mean and percepctr_var: Mean and variance of perceptual spectral contrast.
- 11.tempo: Estimated tempo of the song in beats per minute (BPM).
- 12.mfcc1_mean to mfcc20_mean: Mean values of the first 20 Mel-frequency cepstral coefficients (MFCCs).
- 13.mfcc1_var to mfcc20_var: Variance of the first 20 MFCCs.
- 14.label: The genre label for each audio sample (e.g., blues, classical, jazz).

```
#import the dataset
dt=pd.read_csv("/content/features_3_sec.csv")
dt

{"type": "dataframe", "variable_name": "dt"}

#Total number of Rows and Colomns
dt.shape
```

```
(9990, 60)
```

```
#Read the first 5 rows of dataset
```

```
dt.head()
```

```
{"type": "dataframe", "variable_name": "dt"}
```

```
#Understanding the central Tendency
```

```
dt.describe()
```

```
{"type": "dataframe"}
```

```
dt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 9990 entries, 0 to 9989
```

```
Data columns (total 60 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|-------------------------|----------------|---------|
| 0 | filename | 9990 non-null | object |
| 1 | length | 9990 non-null | int64 |
| 2 | chroma_stft_mean | 9990 non-null | float64 |
| 3 | chroma_stft_var | 9990 non-null | float64 |
| 4 | rms_mean | 9990 non-null | float64 |
| 5 | rms_var | 9990 non-null | float64 |
| 6 | spectral_centroid_mean | 9990 non-null | float64 |
| 7 | spectral_centroid_var | 9990 non-null | float64 |
| 8 | spectral_bandwidth_mean | 9990 non-null | float64 |
| 9 | spectral_bandwidth_var | 9990 non-null | float64 |
| 10 | rolloff_mean | 9990 non-null | float64 |
| 11 | rolloff_var | 9990 non-null | float64 |
| 12 | zero_crossing_rate_mean | 9990 non-null | float64 |
| 13 | zero_crossing_rate_var | 9990 non-null | float64 |
| 14 | harmony_mean | 9990 non-null | float64 |
| 15 | harmony_var | 9990 non-null | float64 |
| 16 | perceptr_mean | 9990 non-null | float64 |
| 17 | perceptr_var | 9990 non-null | float64 |
| 18 | tempo | 9990 non-null | float64 |
| 19 | mfcc1_mean | 9990 non-null | float64 |
| 20 | mfcc1_var | 9990 non-null | float64 |
| 21 | mfcc2_mean | 9990 non-null | float64 |
| 22 | mfcc2_var | 9990 non-null | float64 |
| 23 | mfcc3_mean | 9990 non-null | float64 |
| 24 | mfcc3_var | 9990 non-null | float64 |
| 25 | mfcc4_mean | 9990 non-null | float64 |
| 26 | mfcc4_var | 9990 non-null | float64 |
| 27 | mfcc5_mean | 9990 non-null | float64 |
| 28 | mfcc5_var | 9990 non-null | float64 |
| 29 | mfcc6_mean | 9990 non-null | float64 |
| 30 | mfcc6_var | 9990 non-null | float64 |
| 31 | mfcc7_mean | 9990 non-null | float64 |

| | | | | |
|----|-------------|------|----------|---------|
| 32 | mfcc7_var | 9990 | non-null | float64 |
| 33 | mfcc8_mean | 9990 | non-null | float64 |
| 34 | mfcc8_var | 9990 | non-null | float64 |
| 35 | mfcc9_mean | 9990 | non-null | float64 |
| 36 | mfcc9_var | 9990 | non-null | float64 |
| 37 | mfcc10_mean | 9990 | non-null | float64 |
| 38 | mfcc10_var | 9990 | non-null | float64 |
| 39 | mfcc11_mean | 9990 | non-null | float64 |
| 40 | mfcc11_var | 9990 | non-null | float64 |
| 41 | mfcc12_mean | 9990 | non-null | float64 |
| 42 | mfcc12_var | 9990 | non-null | float64 |
| 43 | mfcc13_mean | 9990 | non-null | float64 |
| 44 | mfcc13_var | 9990 | non-null | float64 |
| 45 | mfcc14_mean | 9990 | non-null | float64 |
| 46 | mfcc14_var | 9990 | non-null | float64 |
| 47 | mfcc15_mean | 9990 | non-null | float64 |
| 48 | mfcc15_var | 9990 | non-null | float64 |
| 49 | mfcc16_mean | 9990 | non-null | float64 |
| 50 | mfcc16_var | 9990 | non-null | float64 |
| 51 | mfcc17_mean | 9990 | non-null | float64 |
| 52 | mfcc17_var | 9990 | non-null | float64 |
| 53 | mfcc18_mean | 9990 | non-null | float64 |
| 54 | mfcc18_var | 9990 | non-null | float64 |
| 55 | mfcc19_mean | 9990 | non-null | float64 |
| 56 | mfcc19_var | 9990 | non-null | float64 |
| 57 | mfcc20_mean | 9990 | non-null | float64 |
| 58 | mfcc20_var | 9990 | non-null | float64 |
| 59 | label | 9990 | non-null | object |

dtypes: float64(57), int64(1), object(2)

memory usage: 4.6+ MB

#check if any null values present

dt.isnull()

{"type": "dataframe"}

#total no. of null values in each column

dt.isnull().sum()

| | |
|-------------------------|---|
| filename | 0 |
| length | 0 |
| chroma_stft_mean | 0 |
| chroma_stft_var | 0 |
| rms_mean | 0 |
| rms_var | 0 |
| spectral_centroid_mean | 0 |
| spectral_centroid_var | 0 |
| spectral_bandwidth_mean | 0 |
| spectral_bandwidth_var | 0 |
| rolloff_mean | 0 |

| | |
|-------------------------|---|
| rolloff_var | 0 |
| zero_crossing_rate_mean | 0 |
| zero_crossing_rate_var | 0 |
| harmony_mean | 0 |
| harmony_var | 0 |
| perceptr_mean | 0 |
| perceptr_var | 0 |
| tempo | 0 |
| mfcc1_mean | 0 |
| mfcc1_var | 0 |
| mfcc2_mean | 0 |
| mfcc2_var | 0 |
| mfcc3_mean | 0 |
| mfcc3_var | 0 |
| mfcc4_mean | 0 |
| mfcc4_var | 0 |
| mfcc5_mean | 0 |
| mfcc5_var | 0 |
| mfcc6_mean | 0 |
| mfcc6_var | 0 |
| mfcc7_mean | 0 |
| mfcc7_var | 0 |
| mfcc8_mean | 0 |
| mfcc8_var | 0 |
| mfcc9_mean | 0 |
| mfcc9_var | 0 |
| mfcc10_mean | 0 |
| mfcc10_var | 0 |
| mfcc11_mean | 0 |
| mfcc11_var | 0 |
| mfcc12_mean | 0 |
| mfcc12_var | 0 |
| mfcc13_mean | 0 |
| mfcc13_var | 0 |
| mfcc14_mean | 0 |
| mfcc14_var | 0 |
| mfcc15_mean | 0 |
| mfcc15_var | 0 |
| mfcc16_mean | 0 |
| mfcc16_var | 0 |
| mfcc17_mean | 0 |
| mfcc17_var | 0 |
| mfcc18_mean | 0 |
| mfcc18_var | 0 |
| mfcc19_mean | 0 |
| mfcc19_var | 0 |
| mfcc20_mean | 0 |
| mfcc20_var | 0 |

```
label          0
dtype: int64
```

```
#Getting more info about dataset
dt.info
```

```
<bound method DataFrame.info of
chroma_stft_mean  chroma_stft_var  rms_mean  \
0      blues.00000.0.wav    66149      0.335406      0.091048
0.130405
1      blues.00000.1.wav    66149      0.343065      0.086147
0.112699
2      blues.00000.2.wav    66149      0.346815      0.092243
0.132003
3      blues.00000.3.wav    66149      0.363639      0.086856
0.132565
4      blues.00000.4.wav    66149      0.335579      0.088129
0.143289
...
...
9985   rock.00099.5.wav    66149      0.349126      0.080515
0.050019
9986   rock.00099.6.wav    66149      0.372564      0.082626
0.057897
9987   rock.00099.7.wav    66149      0.347481      0.089019
0.052403
9988   rock.00099.8.wav    66149      0.387527      0.084815
0.066430
9989   rock.00099.9.wav    66149      0.369293      0.086759
0.050524
```

```
      rms_var  spectral_centroid_mean  spectral_centroid_var  \
0      0.003521      1773.065032      167541.630869
1      0.001450      1816.693777      90525.690866
2      0.004620      1788.539719      111407.437613
3      0.002448      1655.289045      111952.284517
4      0.001701      1630.656199      79667.267654
...
...
9985   0.000097      1499.083005      164266.886443
9986   0.000088      1847.965128      281054.935973
9987   0.000701      1346.157659      662956.246325
9988   0.000320      2084.515327      203891.039161
9989   0.000067      1634.330126      411429.169769
```

```
      spectral_bandwidth_mean  spectral_bandwidth_var  ...  mfcc16_var
\
0      1972.744388      117335.771563  ...  39.687145
1      2010.051501      65671.875673  ...  64.748276
```

| | | | | |
|------|-------------|---------------|-----|-----------|
| 2 | 2084.565132 | 75124.921716 | ... | 67.336563 |
| 3 | 1960.039988 | 82913.639269 | ... | 47.739452 |
| 4 | 1948.503884 | 60204.020268 | ... | 30.336359 |
| ... | ... | ... | ... | ... |
| 9985 | 1718.707215 | 85931.574523 | ... | 42.485981 |
| 9986 | 1906.468492 | 99727.037054 | ... | 32.415203 |
| 9987 | 1561.859087 | 138762.841945 | ... | 78.228149 |
| 9988 | 2018.366254 | 22860.992562 | ... | 28.323744 |
| 9989 | 1867.422378 | 119722.211518 | ... | 38.801735 |

| | mfcc17_mean | mfcc17_var | mfcc18_mean | mfcc18_var | mfcc19_mean \ |
|------|-------------|------------|-------------|------------|---------------|
| 0 | -3.241280 | 36.488243 | 0.722209 | 38.099152 | -5.050335 |
| 1 | -6.055294 | 40.677654 | 0.159015 | 51.264091 | -2.837699 |
| 2 | -1.768610 | 28.348579 | 2.378768 | 45.717648 | -1.938424 |
| 3 | -3.841155 | 28.337118 | 1.218588 | 34.770935 | -3.580352 |
| 4 | 0.664582 | 45.880913 | 1.689446 | 51.363583 | -3.392489 |
| ... | ... | ... | ... | ... | ... |
| 9985 | -9.094270 | 38.326839 | -4.246976 | 31.049839 | -5.625813 |
| 9986 | -12.375726 | 66.418587 | -3.081278 | 54.414265 | -11.960546 |
| 9987 | -2.524483 | 21.778994 | 4.809936 | 25.980829 | 1.775686 |
| 9988 | -5.363541 | 17.209942 | 6.462601 | 21.442928 | 2.354765 |
| 9989 | -11.598399 | 58.983097 | -0.178517 | 55.761299 | -6.903252 |

| | mfcc19_var | mfcc20_mean | mfcc20_var | label |
|------|------------|-------------|------------|-------|
| 0 | 33.618073 | -0.243027 | 43.771767 | blues |
| 1 | 97.030830 | 5.784063 | 59.943081 | blues |
| 2 | 53.050835 | 2.517375 | 33.105122 | blues |
| 3 | 50.836224 | 3.630866 | 32.023678 | blues |
| 4 | 26.738789 | 0.536961 | 29.146694 | blues |
| ... | ... | ... | ... | ... |
| 9985 | 48.804092 | 1.818823 | 38.966969 | rock |
| 9986 | 63.452255 | 0.428857 | 18.697033 | rock |
| 9987 | 48.582378 | -0.299545 | 41.586990 | rock |
| 9988 | 24.843613 | 0.675824 | 12.787750 | rock |
| 9989 | 39.485901 | -3.412534 | 31.727489 | rock |

```
[9990 rows x 60 columns]>
```

```
#Getting no of unique values in length column
dt["length"].nunique()
```

```
1
```


```

# Let's visualize all the music genres plotted according to two random
# features. We use the seaborn library to make a scatterplot of the two
# random features.
feature_names = dt.keys()[::-1]
x_name = random.choice(feature_names)
y_name = random.choice(feature_names)

while x_name == y_name:
    y_name = random.choice(feature_names)

sns.relplot(x = x_name, y = y_name, hue = "label", data = dt);

```



```

#make a copy of dataset to work on
ndt=dt.copy()

#here we will work on ndt it will be copy dataset .
#here remove all non float values except genre so our length and name
#id column will get dropped.
non_floats = []
for col in ndt.iloc[:,::-1]:
    if ndt[col].dtypes != "float64":

```

```

        non_floats.append(col)
ndt = ndt.drop(columns=non_floats)

ndt

{"type": "dataframe", "variable_name": "ndt"}

```

What are the Most Common Genres in the Dataset?

```

ndt["label"].value_counts()

label
blues      1000
jazz       1000
metal      1000
pop        1000
reggae     1000
disco      999
classical  998
hiphop     998
rock       998
country    997
Name: count, dtype: int64

#Categorisation of data
#feature catergorization
#determine x and y
x=ndt.iloc[:,0:57].values    #:, means all values
y=ndt.iloc[:,57].values

#Label encoding
!pip install --user scikit-learn
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y=le.fit_transform(y)

Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.10/dist-packages (1.5.2)
Requirement already satisfied: numpy>=1.19.5 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.26.4)
Requirement already satisfied: scipy>=1.6.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)

#building machine learning models
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import

```



```

accuracy_score,precision_score,recall_score,f1_score

#model selection
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB

#splittinX
x_train,x_test,y_train,y_test =
train_test_split(x,y,test_size=0.3,random_state=0)

```

Which Machine Learning Model gives the highest accuracy, Precision , F1 Score, Recall?

```

#Accuracy, Precision, F1 score,Recall,Confusion Matrix using Guassian
Naive Bayes Algorithm
accuracy_list=[]
model_name=[]
gaussian = GaussianNB()
gaussian.fit(x_train,y_train)
Y_pred = gaussian.predict(x_test)
accuracy_nb=round(accuracy_score(y_test,Y_pred)* 100, 2 )

cm = confusion_matrix(y_test, Y_pred)
accuracy = accuracy_score(y_test,Y_pred)
accuracy_list.append(accuracy)
model_name.append("GNB")
print("accuracy_Naive Bayes: %.3f" %accuracy)

accuracy_Naive Bayes: 0.425

pre_nb=precision_score(y_test,Y_pred,average='weighted')
print("precision_Naive Bayes: %.3f" %pre_nb)
recall_score_nb=recall_score(y_test,Y_pred,average='weighted')
print("recall_Naive Bayes: %.3f" %recall_score_nb)
f1_score_nb=f1_score(y_test,Y_pred,average='weighted')
print("f1_score_Naive Bayes: %.3f" %f1_score_nb)
confusion_matrix_nb=confusion_matrix(y_test,Y_pred)
print("confusion_matrix_Naive Bayes:\n", confusion_matrix_nb)

precision_Naive Bayes: 0.439
recall_Naive Bayes: 0.425
f1_score_Naive Bayes: 0.402
confusion_matrix_Naive Bayes:
[[ 64  18  39   5   3  30 101   0  21   6]
 [  1 262   3   1   0   7  12   1   3   5]
 [ 19  16 103  47   5  11  71   1  24  16]
 [  9   2  10 120  11   3  92   8  21  14]
 [  7   0  27  57  75   1  55  25  47   7]
 [ 24  63  13  33   0  68  52   8   7  27]
 [  1   2   2  16   8   2 274   1   2   4]
 [  1   2   6  76  11   3  19 140  23   6]

```

```
[ 32  1 31 36 31  2  8 27 142  8]
[  3 12 38 42 10  6 134  6 21 27]]
```

```
from sklearn import preprocessing
from sklearn.model_selection import cross_val_score
from sklearn import svm
```

```
from sklearn.linear_model import LogisticRegression
```

#Accuracy, Precision, F1 score, Recall, Confusion Matrix using Logistic Regression Algorithm

```
model1=LogisticRegression(max_iter=500,random_state=70)
model1.fit(x_train,y_train)
y_pred=model1.predict(x_test)
accuracy=accuracy_score(y_test,y_pred)
print(accuracy)
accuracy_list.append(accuracy)
model_name.append("LR")
```

```
0.30430430430430433
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
pre_lr=precision_score(y_test,y_pred,average='weighted')
print("precision_Logistic Regression: %.3f" %pre_lr)
recall_score_lr=recall_score(y_test,y_pred,average='weighted')
print("recall_Logistic Regression: %.3f" %recall_score_lr)
f1_score_lr=f1_score(y_test,y_pred,average='weighted')
print("f1_score_Logistic Regression: %.3f" %f1_score_lr)
confusion_matrix_lr=confusion_matrix(y_test,y_pred)
print("confusion_matrix_Logistic Regression:\n", confusion_matrix_lr)
```

```
precision_Logistic Regression: 0.296
recall_Logistic Regression: 0.304
f1_score_Logistic Regression: 0.269
confusion_matrix_Logistic Regression:
[[ 50 39  4 34  6 38 73 16 26  1]
 [  8 141  4  2  0  8 128  0  4  0]
 [ 23 35 37 47  4 48 42 40 28  9]
```

```
[ 15  4 14 74 10 16 71 64 21 1]
[ 11  7 15 55 25  6 47 87 46 2]
[ 41 48 10 18  2 73 64 32  1 6]
[  1 11  7 45  1  6 235  4  2 0]
[  4  5  6 34 13 23 22 152 25 3]
[ 40 11 19 32 27  9 13 46 118 3]
[ 14 18 11 62  4 28 83 53 19 7]]
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
#Accuracy, Precision, F1 score,Recall,Confusion Matrix using
DecisionTree Classifier Algorithm
```

```
model2=DecisionTreeClassifier(random_state=42)
```

```
model2.fit(x_train,y_train)
```

```
y_dpred=model2.predict(x_test)
```

```
accuracy1=accuracy_score(y_test,y_dpred)
```

```
print(accuracy1)
```

```
accuracy_list.append(accuracy1)
```

```
model_name.append("DTC")
```

```
0.6489823156489823
```

```
pre_dtc=precision_score(y_test,y_dpred,average='weighted')
```

```
print("precision_Decision Tree Classifier: %.3f" %pre_dtc)
```

```
recall_score_dtc=recall_score(y_test,y_dpred,average='weighted')
```

```
print("recall_Decision Tree Classifier: %.3f" %recall_score_dtc)
```

```
f1_score_dtc=f1_score(y_test,y_dpred,average='weighted')
```

```
print("f1_score_Decision Tree Classifier: %.3f" %f1_score_dtc)
```

```
confusion_matrix_dtc=confusion_matrix(y_test,y_dpred)
```

```
print("confusion_matrix_Decision Tree Classifier:\n",
confusion_matrix_dtc)
```

```
precision_Decision Tree Classifier: 0.650
```

```
recall_Decision Tree Classifier: 0.649
```

```
f1_score_Decision Tree Classifier: 0.649
```

```
confusion_matrix_Decision Tree Classifier:
```

```
[[175  1 25 11  5 11 17  1 17 24]
[  1 263  7  0  1 17  0  0  1  5]
[ 32  6 164 18  6 21  6  7 13 40]
[  9  4 10 163 28  5  7 11 24 29]
[  7  1  7 25 198  1  5 24 25  8]
[ 14 28 23  4  2 194  3  3  8 16]
[ 11  0  3  6 14  6 234  1  6 31]
[  2  2 11 21 17  7  0 198 20  9]
[ 12  2 14 21 21  7  4 13 214 10]
[ 18  6 28 32  7 18 27  6 15 142]]
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
#Accuracy, Precision, F1 score,Recall,Confusion Matrix using
RandomForestClassifier Algorithm
```

```

model3=RandomForestClassifier(n_estimators=100,random_state=1)
model3.fit(x_train,y_train)
y_rpred=model3.predict(x_test)
accuracy2=accuracy_score(y_test,y_rpred)
accuracy_list.append(accuracy2)
model_name.append("RFC")
print(accuracy2)

```

0.8638638638638638

```

pre_rfc=precision_score(y_test,y_rpred,average='weighted')
print("precision_Random Forest Classifier: %.3f" %pre_rfc)
recall_score_rfc=recall_score(y_test,y_rpred,average='weighted')
print("recall_Random Forest Classifier: %.3f" %recall_score_rfc)
f1_score_rfc=f1_score(y_test,y_rpred,average='weighted')
print("f1_score_Random Forest Classifier: %.3f" %f1_score_rfc)
confusion_matrix_rfc=confusion_matrix(y_test,y_rpred)
print("confusion_matrix_Random Forest Classifier:\n",
confusion_matrix_rfc)

```

```

precision_Random Forest Classifier: 0.864
recall_Random Forest Classifier: 0.864
f1_score_Random Forest Classifier: 0.862
confusion_matrix_Random Forest Classifier:
[[265  1  3  4  1  7  4  0  2  0]
 [  0 286  3  0  0  6  0  0  0  0]
 [ 16  1 267  3  1 10  3  1  7  4]
 [  4  4  5 248  9  1  3  2  5  9]
 [  1  1  3  9 253  0  5 17  8  4]
 [  5 16 12  3  0 258  0  1  0  0]
 [  2  1  1  2  5  2 286  0  1 12]
 [  0  2  9  9  6  1  0 250  7  3]
 [  5  2 12  5  6  2  0  9 275  2]
 [ 14  3 17 21  5 15 17  2  4 201]]

```

```

from sklearn.ensemble import AdaBoostClassifier

```

*#Accuracy, Precision, F1 score,Recall,Confusion Matrix using
AdaboostClassifier Algorithm*

```

model_ada=AdaBoostClassifier(model2,learning_rate=1.0,random_state=42)
model_ada.fit(x_train,y_train)
y_apred=model_ada.predict(x_test)
accu=accuracy_score(y_test,y_apred)
accuracy_list.append(accu)
model_name.append("ABC")
print(accu)

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME

```
algorithm to circumvent this warning.  
warnings.warn(  

```

```
0.6866866866866866
```

```
pre_dtc=precision_score(y_test,y_apred,average='weighted')  
print("precision_AdaBoost Classifier: %.3f" %pre_dtc)  
recall_score_dtc=recall_score(y_test,y_apred,average='weighted')  
print("recall_AdaBoost Classifier: %.3f" %recall_score_dtc)  
f1_score_dtc=f1_score(y_test,y_apred,average='weighted')  
print("f1_score_AdaBoost Classifier: %.3f" %f1_score_dtc)  
confusion_matrix_dtc=confusion_matrix(y_test,y_apred)  
print("confusion_matrix_AdaBoost Classifier:\n", confusion_matrix_dtc)
```

```
precision_AdaBoost Classifier: 0.688  
recall_AdaBoost Classifier: 0.687  
f1_score_AdaBoost Classifier: 0.687  
confusion_matrix_AdaBoost Classifier:  
[[190  0  27  7  10  17  8  2  13  13]  
 [  0 256  8  0  0  23  2  0  1  5]  
 [ 24  3 185 18  7  26  5  6 14 25]  
 [ 16  3  12 172 16  5 10 10 21 25]  
 [  8  0  9  14 197  2 10 30 22  9]  
 [  8 15  27  5  5 216  1  7  2  9]  
 [  7  0  2  9  7  2 256  0  1 28]  
 [  1  5 14 12 10  7  0 213 14 11]  
 [ 13  2 17 15 26  3  5  9 211 17]  
 [ 21  3 31 20  7 16 22  8  9 162]]
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
#Accuracy, Precision, F1 score,Recall,Confusion Matrix using Gradient  
BoostingClassifier Algorithm  
model_gra=GradientBoostingClassifier(n_estimators=100,max_depth=3,rand  
om_state=42)  
model_gra.fit(x_train,y_train)  
y_gpred=model_gra.predict(x_test)  
accu_g=accuracy_score(y_test,y_gpred)  
print(accu_g)  
accuracy_list.append(accu_g)  
model_name.append("GBC")
```

```
0.8218218218218218
```

```
pre_gbc=precision_score(y_test,y_gpred,average='weighted')  
print("Gradient_Boosting Classifier: %.3f" %pre_gbc)  
recall_score_gbc=recall_score(y_test,y_apred,average='weighted')  
print("Gradient_Boosting Classifier: %.3f" %recall_score_gbc)  
f1_score_gbc=f1_score(y_test,y_apred,average='weighted')  
print("Gradient_Boosting Classifier: %.3f" %f1_score_gbc)  
confusion_matrix_gbc=confusion_matrix(y_test,y_gpred)
```

```
print("confusion_matrix_Gradient_Boosting Classifier:\n",
confusion_matrix_gbc)
```

```
Gradient_Boosting Classifier: 0.822
```

```
Gradient_Boosting Classifier: 0.687
```

```
Gradient_Boosting Classifier: 0.687
```

```
confusion_matrix_Gradient_Boosting Classifier:
```

```
[[243  1  10  8  5  7  4  0  4  5]
 [ 0 275  4  0  0 10  0  0  1  5]
 [ 19  2 244  6  0 13  5  3  9 12]
 [ 3  3  8 232 16  3  3  7  8  7]
 [ 4  1  5  5 242  1  5 17 15  6]
 [ 2 17 19  1  0 254  0  0  0  2]
 [ 5  1  1  8  6  1 269  0  0 21]
 [ 0  2  8 15  8  4  0 238  8  4]
 [ 6  1 12  7  9  0  0  9 262 12]
 [16  4 20 18  5 11 13  4  4 204]]
```

```
#Accuracy, Precision, F1 score,Recall,Confusion Matrix using SVC
Algorithm
```

```
from sklearn.svm import SVC
```

```
from sklearn.model_selection import GridSearchCV
```

```
svm = SVC(kernel='rbf')
```

```
# Hyperparameter tuning
```

```
param_grid = {'C': [0.5,1,10], 'gamma': ['scale', 'auto']}
```

```
grid_search = GridSearchCV(svm, param_grid, cv=5)
```

```
grid_search.fit(x_train, y_train)
```

```
# Evaluate model
```

```
best_model = grid_search.best_estimator_
```

```
y_predsvm = best_model.predict(x_test)
```

```
accuracy = accuracy_score(y_test, y_predsvm)
```

```
acc_svm=accuracy_score(y_test,y_predsvm)
```

```
print("Accuracy:",acc_svm)
```

```
model_name.append("SVM")
```

```
accuracy_list.append(acc_svm)
```

```
Accuracy: 0.314647981314648
```

```
pre_svm=precision_score(y_test,y_predsvm,average='weighted')
```

```
print("Gradient_Boosting Classifier: %.3f" %pre_svm)
```

```
recall_score_svm=recall_score(y_test,y_predsvm,average='weighted')
```

```
print("Gradient_Boosting Classifier: %.3f" %recall_score_svm)
```

```
f1_score_svm=f1_score(y_test,y_predsvm,average='weighted')
```

```
print("Gradient_Boosting Classifier: %.3f" %f1_score_svm)
```

```
confusion_matrix_svm=confusion_matrix(y_test,y_predsvm)
```

```
print("confusion_matrix_Gradient_Boosting Classifier:\n",
```

```
confusion_matrix_svm)
```

#Accuracy, Precision, F1 score, Recall, Confusion Matrix using xgboost Algorithm

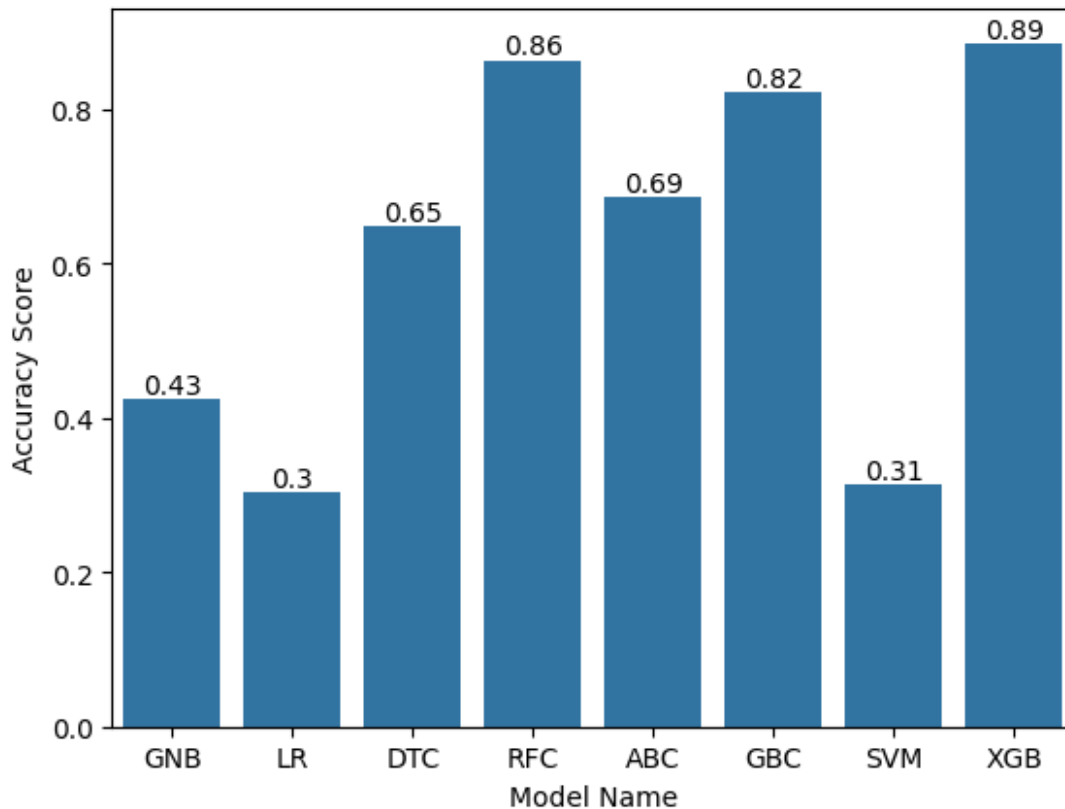
```
import xgboost as xgb
xgb_model=xgb.XGBClassifier()
xgb_model.fit(x_train,y_train)
y_predxgb=xgb_model.predict(x_test)
acc_xgb=accuracy_score(y_test,y_predxgb)
print("Accuracy:",acc_xgb)
model_name.append("XGB")
accuracy_list.append(acc_xgb)

pre_xgb=precision_score(y_test,y_predxgb,average='weighted')
print("xgboost Classifier: %.3f" %pre_xgb)
recall_score_xgb=recall_score(y_test,y_predxgb,average='weighted')
print("xgboost Classifier: %.3f" %recall_score_xgb)
f1_score_xgb=f1_score(y_test,y_predxgb,average='weighted')
print("xgboost Classifier: %.3f" %f1_score_xgb)
confusion_matrix_xgb=confusion_matrix(y_test,y_predxgb)
print("confusion_matrix_xgboost Classifier:\n", confusion_matrix_xgb)
```

```
xgboost Classifier: 0.887
xgboost Classifier: 0.886
xgboost Classifier: 0.886
confusion_matrix_xgboost Classifier:
[[263   1   5   5   2   5   3   0   0   3]
 [  0 280   2   0   0   8   0   0   0   5]
 [ 11   2 275   2   1   8   0   0   4  10]
 [  6   4   6 248   7   1   1   4   6   7]
 [  2   1   4   4 269   1   3   8   7   2]
 [  1  11  12   1   0 268   1   0   0   1]
 [  4   1   1   2   6   1 287   0   0  10]
 [  0   1   6   5   6   2   0 256   6   5]
 [  6   2   6   4  10   1   0   8 277   4]
 [ 10   5  13   9   4   8  10   3   4 233]]
```

#Comparing the accuracy of all the models and determining the model with greatest accuracy

```
new_df=pd.DataFrame({
    "Model Name":model_name,"Accuracy Score":accuracy_list
})
bar_plot=sns.barplot(new_df,x="Model Name",y="Accuracy Score")
for index,row in new_df.iterrows():
    bar_plot.text(index,row["Accuracy Score"],round(row["Accuracy Score"],2),color="black",ha="center",va="bottom")
plt.show()
```



On which top 5 features accuracy is highly dependent?

```
imp_list=xgb_model.feature_importances_
model_imp=pd.Series(imp_list,ndt.columns[:-1])
top_5=model_imp.nlargest(5)
print(top_5)
```


| | |
|-------------------------|----------|
| perceptr_var | 0.072947 |
| spectral_bandwidth_mean | 0.057957 |
| chroma_stft_mean | 0.040678 |
| mfcc4_mean | 0.037323 |
| rms_var | 0.032699 |

```
dtype: float32
```

```
#Comparing the above top 5 feature dependency
feat_name=["perceptr_var","s_band_mean","chr_stft_mean","mfcc4_mean","rms_var"]
new_df2=pd.DataFrame({
    "feat_name":feat_name,"top_5":top_5
})
bar_plot2=sns.barplot(x=feat_name,y=top_5)
plt.xlabel("Feature Name")
plt.ylabel("Dependancy")
plt.title("Feature Dependancy on Model Prediction")
#for index,row in new_df2.iterrows():
```



```
#
bar_plot2.text(index, row["top_5"], round(row["top_5"], 2), color="black",
ha="center", va="top")
plt.show()
```




Are There Any Correlations Between Features?

```
correlation = ndt.iloc[:,0:57].corr()
correlation

{"type": "dataframe", "variable_name": "correlation"}

sns.heatmap(correlation, square=True)

<Axes: >
```



CONCLUSION:

#In this project, we have built a model which predicts genre of the music based on some features of the music.
 #from above analysis we concluded that among naive bayes, logistic regression, decision tree classifier,
 #random forest classifier, adaboost classifier, xgboost classifier
 #, gradient boosting classifier the most accurate model is xgboost classifier with accuracy of almost 89% .
 #Using the model we analysed the top five features on which the accuracy of this model are most dependent.
 #Among these the topmost was perceptr_var followed by
 ms_var, chroma_stft_mean, rms_mean, mfcc4_mean.